

# Cache memories

**Slides by:** Randal E. Bryant and David R. O'Hallaron (CMU)

Presented by Xin Jin and David Hovemeyer for CSF

October 15, 2021

601.229 Computer Systems Fundamentals



# Cache writes and performance

# What about writes?

- ▶ **Multiple copies of data exist:**
  - ▶ L1, L2, L3, Main Memory, Disk
- ▶ **What to do on a write-hit?**
  - ▶ **Write-through** (write immediately to memory)
  - ▶ **Write-back** (defer write to memory until replacement of line)
    - ▶ Need a dirty bit (line different from memory or not)
- ▶ **What to do on a write-miss?**
  - ▶ **Write-allocate** (load into cache, update line in cache)
    - ▶ Good if more writes to the location follow
  - ▶ **No-write-allocate** (writes straight to memory, does not load into cache)
- ▶ **Typical**
  - ▶ Write-through + No-write-allocate
  - ▶ **Write-back + Write-allocate**

# Zoom poll #1!

Consider the following code:

```
for (int i = 0; i < 8; i++) {  
    a[i] = i * 2;  
}
```

Assume that

- ▶ address of a[0] is a multiple of 16
- ▶ cache is cold initially
- ▶ i is a register
- ▶ sizeof(int)=4
- ▶ 16 bytes per block
- ▶ cache is direct-mapped
- ▶ loads and stores always access exactly 4 bytes

If there are **8** stores to *memory*, what cache configuration is likely?

- A. write-allocate + write-through
- B. no-write-allocate + write-through
- C. write-allocate + write-back
- D. no-write-allocate + write-back

# Zoom poll #2!

Consider the following code:

```
for (int i = 0; i < 8; i++) {  
    a[i] = i * 2;  
}
```

Assume that

- ▶ address of a[0] is a multiple of 16
- ▶ cache is cold initially
- ▶ i is a register
- ▶ sizeof(int)=4
- ▶ 16 bytes per block
- ▶ cache is direct-mapped
- ▶ loads and stores always access exactly 4 bytes

If the cache is configured for write-allocate + write-back, how many **loads** from *memory* are there?

- A. 0
- B. 2
- C. 8
- D. 10
- E. 16

# Zoom poll #3!

Consider the following code:

```
for (int i = 0; i < 8; i++) {  
    a[i] = i * 2;  
}
```

Assume that

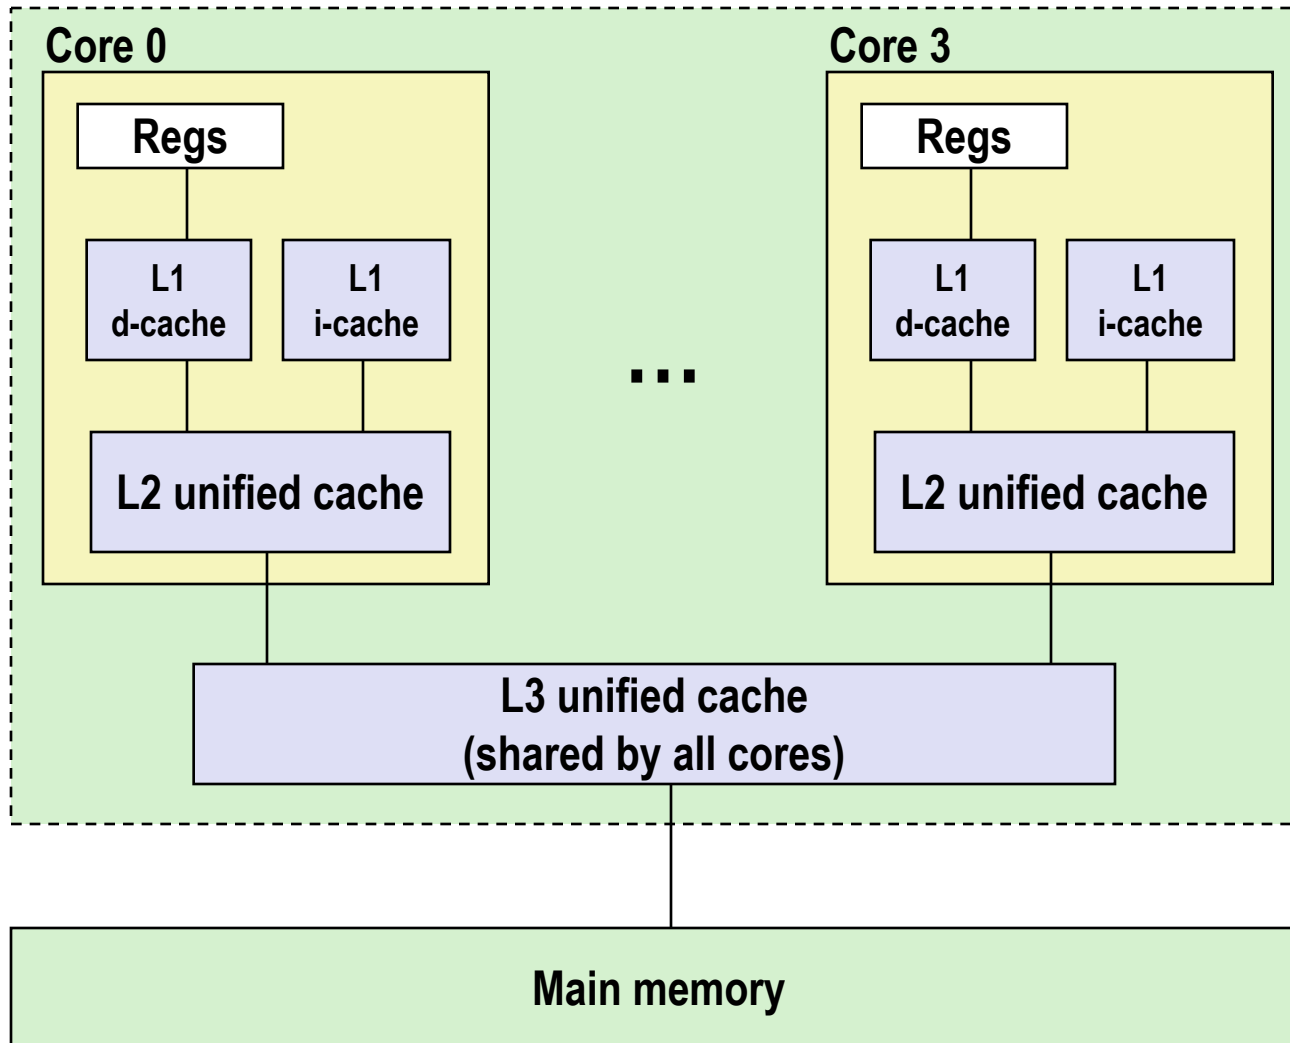
- ▶ address of a[0] is a multiple of 16
- ▶ cache is cold initially
- ▶ i is a register
- ▶ sizeof(int)=4
- ▶ 16 bytes per block
- ▶ cache is direct-mapped
- ▶ loads and stores always access exactly 4 bytes

If the cache is configured for write-allocate + write-back, how many **stores** to *memory* are there?

- A. 0
- B. 2
- C. 8
- D. 10
- E. 16

# Intel Core i7 Cache Hierarchy

## Processor package



**L1 i-cache and d-cache:**  
32 KB, 8-way,  
Access: 4 cycles

**L2 unified cache:**  
256 KB, 8-way,  
Access: 10 cycles

**L3 unified cache:**  
8 MB, 16-way,  
Access: 40-75 cycles

**Block size:** 64 bytes for  
all caches.

# Cache Performance Metrics

## ▶ Miss Rate

- ▶ Fraction of memory references not found in cache (misses / accesses)  
=  $1 - \text{hit rate}$
- ▶ Typical numbers (in percentages):
  - ▶ 3-10% for L1
  - ▶ can be quite small (e.g.,  $< 1\%$ ) for L2, depending on size, etc.

## ▶ Hit Time

- ▶ Time to deliver a line in the cache to the processor
  - ▶ includes time to determine whether the line is in the cache
- ▶ Typical numbers:
  - ▶ 4 clock cycle for L1
  - ▶ 10 clock cycles for L2

## ▶ Miss Penalty

- ▶ Additional time required because of a miss
  - ▶ typically 50-200 cycles for main memory (Trend: increasing!)



# Let's think about those numbers

- ▶ **Huge difference between a hit and a miss**
  - ▶ Could be 100x, if just L1 and main memory
- ▶ **Would you believe 99% hits is twice as good as 97%?**
  - ▶ Consider:  
cache hit time of 1 cycle  
miss penalty of 100 cycles
  - ▶ Average access time:  
97% hits:  $1 \text{ cycle} + 0.03 * 100 \text{ cycles} = 4 \text{ cycles}$   
99% hits:  $1 \text{ cycle} + 0.01 * 100 \text{ cycles} = 2 \text{ cycles}$
- ▶ **This is why “miss rate” is used instead of “hit rate”**

# Writing cache-friendly code

# Writing Cache Friendly Code

- ▶ **Make the common case go fast**
  - ▶ Focus on the inner loops of the core functions
- ▶ **Minimize the misses in the inner loops**
  - ▶ Repeated references to variables are good (**temporal locality**)
  - ▶ Stride-1 reference patterns are good (**spatial locality**)

**Key idea: Our qualitative notion of locality is quantified through our understanding of cache memories**

# Matrix Multiplication Example

## ► Description:

- Multiply N x N matrices
- Matrix elements are doubles (8 bytes)
- $O(N^3)$  total operations
- N reads per source element
- N values summed per destination
  - but may be able to hold in register

```
/* ijk */
for (i=0; i<n; i++) {
    for (j=0; j<n; j++) {
        sum = 0.0;
        for (k=0; k<n; k++)
            sum += a[i][k] * b[k][j];
        c[i][j] = sum;
    }
}
```

*Variable sum  
held in register*

*matmult/mm.c*

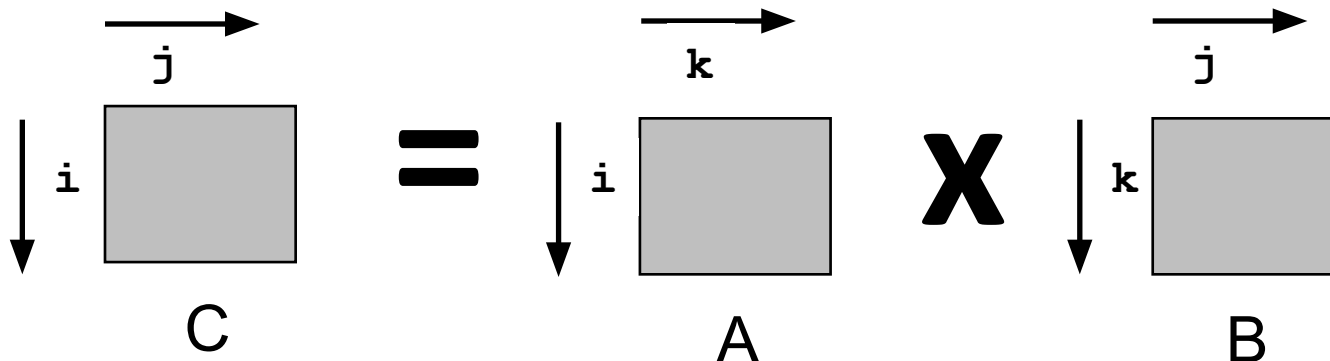
# Miss Rate Analysis for Matrix Multiply

- ▶ **Assume:**

- ▶ Block size =  $32B$  (big enough for four doubles)
- ▶ Matrix dimension ( $N$ ) is very large
  - ▶ Approximate  $1/N$  as  $0.0$
- ▶ Cache is not even big enough to hold multiple rows

- ▶ **Analysis Method:**

- ▶ Look at access pattern of inner loop



# Layout of C Arrays in Memory (review)

- ▶ **C arrays allocated in row-major order**
  - ▶ each row in contiguous memory locations
- ▶ **Stepping through columns in one row:**
  - ▶ 

```
for (i = 0; i < N; i++)  
    sum += a[0][i];
```
  - ▶ accesses successive elements
  - ▶ if block size (B) > sizeof(a<sub>ij</sub>) bytes, exploit spatial locality
    - ▶ miss rate = sizeof(a<sub>ij</sub>) / B
- ▶ **Stepping through rows in one column:**
  - ▶ 


```
for (i = 0; i < n; i++)  
    sum += a[i][0];
```
  - ▶ accesses distant elements
  - ▶ no spatial locality!
    - ▶ miss rate = 1 (i.e. 100%)

# Matrix Multiplication (ijk)

```
/* ijk */  
for (i=0; i<n; i++) {  
    for (j=0; j<n; j++) {  
        sum = 0.0;  
        for (k=0; k<n; k++)  
            sum += a[i][k] * b[k][j];  
        c[i][j] = sum;  
    }  
}
```

*matmult/mm.c*

Inner loop:



Misses per inner loop iteration:


<u>A</u>	<u>B</u>	<u>C</u>
0.25	1.0	0.0

# Matrix Multiplication (jik)

```
/* jik */
for (j=0; j<n; j++) {
    for (i=0; i<n; i++) {
        sum = 0.0;
        for (k=0; k<n; k++)
            sum += a[i][k] * b[k][j];
        c[i][j] = sum
    }
}
```

*matmult/mm.c*

Inner loop:



Misses per inner loop iteration:


<u>A</u>	<u>B</u>	<u>C</u>
0.25	1.0	0.0



# Matrix Multiplication (kij)

```
/* kij */
for (k=0; k<n; k++) {
    for (i=0; i<n; i++) {
        r = a[i][k];
        for (j=0; j<n; j++)
            c[i][j] += r * b[k][j];
    }
}
```

*matmult/mm.c*




Misses per inner loop iteration:

<u>A</u>	<u>B</u>	<u>C</u>
0.0	0.25	0.25

# Matrix Multiplication (ikj)

```
/* ikj */  
for (i=0; i<n; i++) {  
    for (k=0; k<n; k++) {  
        r = a[i][k];  
        for (j=0; j<n; j++)  
            c[i][j] += r * b[k][j];  
    }  
}
```

*matmult/mm.c*



Misses per inner loop iteration:

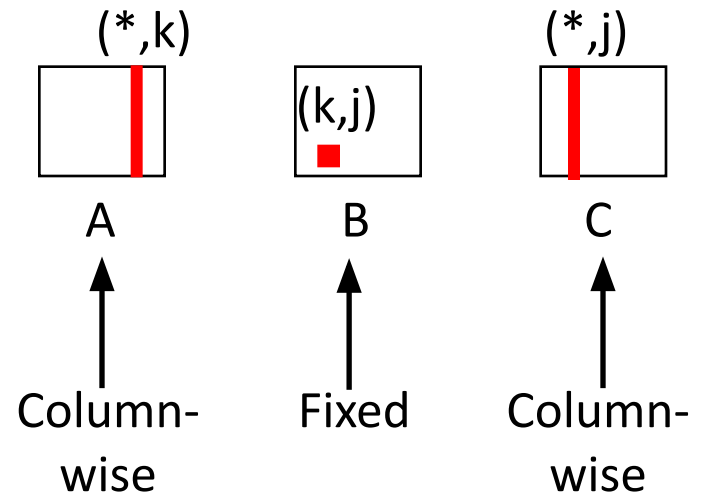
<u>A</u>	<u>B</u>	<u>C</u>
0.0	0.25	0.25

# Matrix Multiplication (jki)

```
/* jki */
for (j=0; j<n; j++) {
    for (k=0; k<n; k++) {
        r = b[k][j];
        for (i=0; i<n; i++)
            c[i][j] += a[i][k] * r;
    }
}
```

*matmult/mm.c*

Inner loop:



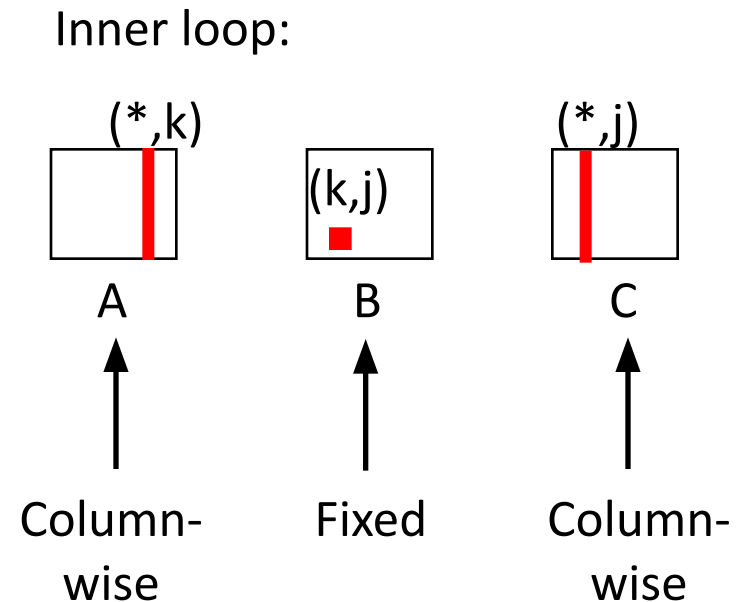
Misses per inner loop iteration:

<u>A</u>	<u>B</u>	<u>C</u>
1.0	0.0	1.0

# Matrix Multiplication (kji)

```
/* kji */  
for (k=0; k<n; k++) {  
    for (j=0; j<n; j++) {  
        r = b[k][j];  
        for (i=0; i<n; i++)  
            c[i][j] += a[i][k] * r;  
    }  
}
```

*matmult/mm.c*



Misses per inner loop iteration:

<u>A</u>	<u>B</u>	<u>C</u>
1.0	0.0	1.0

# Summary of Matrix Multiplication

```
for (i=0; i<n; i++) {  
    for (j=0; j<n; j++) {  
        sum = 0.0;  
        for (k=0; k<n; k++)  
            sum += a[i][k] * b[k][j];  
        c[i][j] = sum;  
    }  
}
```

**ijk (& jik):**

- 2 loads, 0 stores
- misses/iter = **1.25**

```
for (k=0; k<n; k++) {  
    for (i=0; i<n; i++) {  
        r = a[i][k];  
        for (j=0; j<n; j++)  
            c[i][j] += r * b[k][j];  
    }  
}
```

**kij (& ikj):**

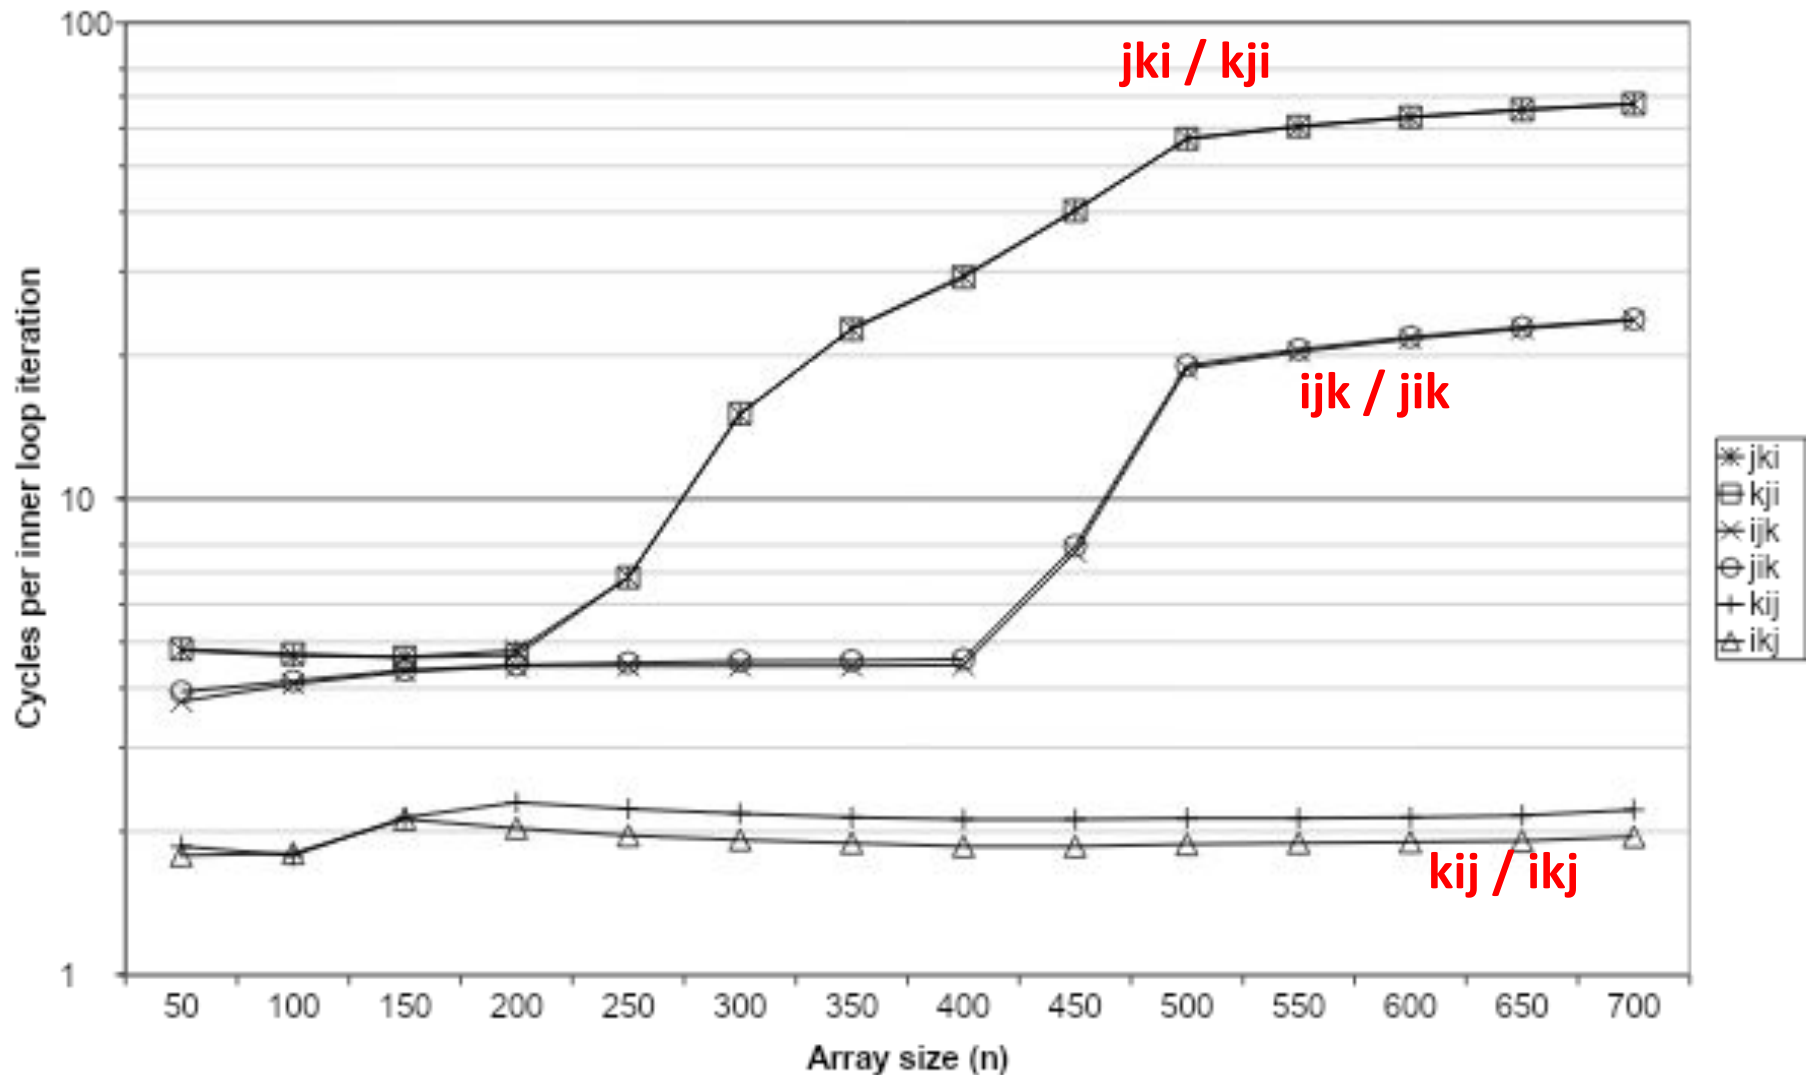
- 2 loads, 1 store
- misses/iter = **0.5**

```
for (j=0; j<n; j++) {  
    for (k=0; k<n; k++) {  
        r = b[k][j];  
        for (i=0; i<n; i++)  
            c[i][j] += a[i][k] * r;  
    }  
}
```

**jki (& kji):**

- 2 loads, 1 store
- misses/iter = **2.0**


# Core i7 Matrix Multiply Performance



# **Use blocking to improve temporal locality**

# Example: Matrix Multiplication

```
c = (double *) calloc(sizeof(double), n*n);  
  
/* Multiply n x n matrices a and b */  
void mmm(double *a, double *b, double *c, int n) {  
    int i, j, k;  
    for (i = 0; i < n; i++)  
        for (j = 0; j < n; j++)  
            for (k = 0; k < n; k++)  
                c[i*n + j] += a[i*n + k] * b[k*n + j];  
}
```





# Cache Miss Analysis


- ▶ **Assume:**

- ▶ Matrix elements are doubles
- ▶ Cache block = 8 doubles
- ▶ Cache size  $C \ll n$  (much smaller than  $n$ )

- ▶ **First iteration:**

- ▶  $n/8 + n = 9n/8$  misses

- ▶ Afterwards **in cache:**  
(schematic)




# Cache Miss Analysis

- ▶ **Assume:**

- ▶ Matrix elements are doubles
- ▶ Cache block = 8 doubles
- ▶ Cache size  $C \ll n$  (much smaller than  $n$ )

- ▶ **Second iteration:**

- ▶ Again:  
 $n/8 + n = 9n/8$  misses



- ▶ **Total misses:**


- ▶  $9n/8 * n^2 = (9/8) * n^3$

# Blocked Matrix Multiplication

```
c = (double *) calloc(sizeof(double), n*n);


/* Multiply n x n matrices a and b */
void mmm(double *a, double *b, double *c, int n) {
    int i, j, k;
    for (i = 0; i < n; i+=B)
        for (j = 0; j < n; j+=B)
            for (k = 0; k < n; k+=B)
                /* B x B mini matrix multiplications */
                for (i1 = i; i1 < i+B; i++)
                    for (j1 = j; j1 < j+B; j++)
                        for (k1 = k; k1 < k+B; k++)
                            c[i1*n+j1] += a[i1*n + k1]*b[k1*n + j1];
}
```

*matmult/bmm.c*



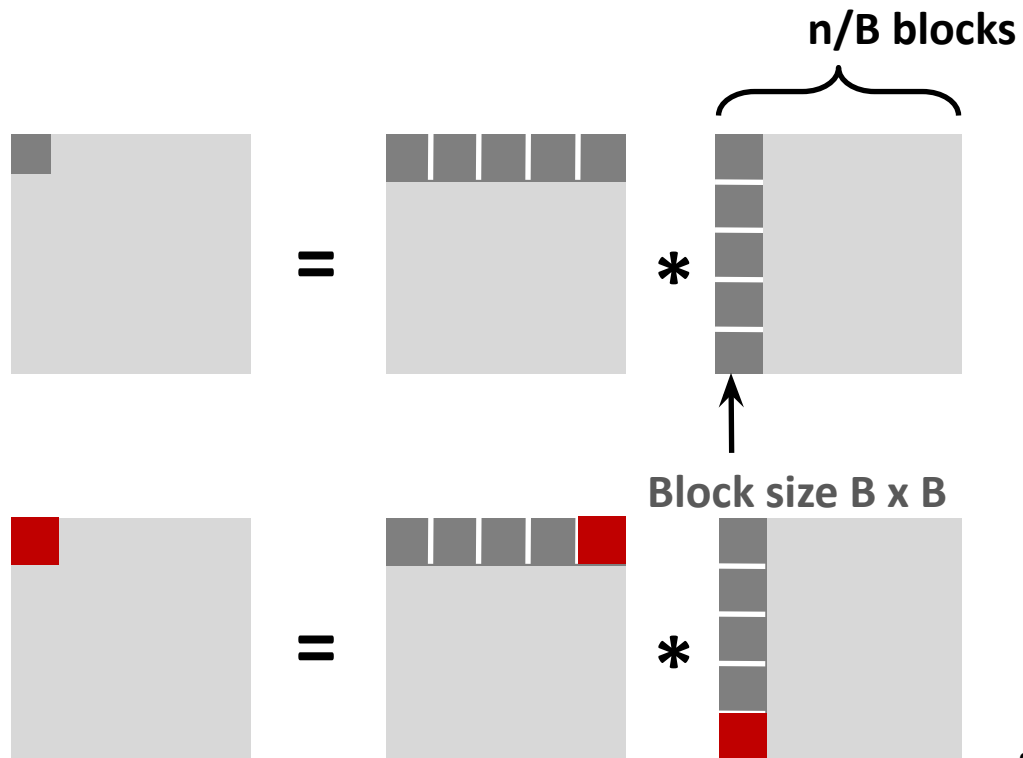
# Cache Miss Analysis

## ► Assume:

- Cache block = 8 doubles
- Cache size  $C \ll n$  (much smaller than  $n$ )
- Three blocks  fit into cache:  $3B^2 < C$


## ► First (block) iteration:

- $B^2/8$  misses for each block
- $2n/B * B^2/8 = nB/4$   
(omitting matrix  $c$ )



# Cache Miss Analysis

## ► Assume:


- Cache block = 8 doubles
- Cache size  $C \ll n$  (much smaller than  $n$ )
- Three blocks  fit into cache:  $3B^2 < C$

## ► Second (block) iteration:

- Same as first iteration
- $2n/B * B^2/8 = nB/4$

## ► Total misses:

- $nB/4 * (n/B)^2 = n^3/(4B)$



# Blocking Summary

- ▶ **No blocking:**  $(9/8) * n^3$
- ▶ **Blocking:**  $1/(4B) * n^3$
- ▶ **Suggest largest possible block size B, but limit  $3B^2 < C$ !**
- ▶ **Reason for dramatic difference:**
  - ▶ Matrix multiplication has inherent temporal locality:
    - ▶ Input data:  $3n^2$ , computation  $2n^3$
    - ▶ Every array elements used  $O(n)$  times!
  - ▶ But program has to be written properly

# Cache Summary

- ▶ **Cache memories can have significant performance impact**
- ▶ **You can write your programs to exploit this!**
  - ▶ Focus on the inner loops, where bulk of computations and memory accesses occur.
  - ▶ Try to maximize spatial locality by reading data objects with sequentially with stride 1.
  - ▶ Try to maximize temporal locality by using a data object as often as possible once it's read from memory.